

```
# Importing the libraries
import pandas as pd
import torch
import torch.nn as nn
import torch.optim as optim
import time
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.impute import SimpleImputer
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

data = pd.read_csv('bank_data.csv')
data.head()
```

Unnamed: 0	id	dependents	education	self_employed
0	1	2	Graduate	No
1	2	0	Not Graduate	Yes
2	3	3	Graduate	No
3	4	3	Graduate	No
4	5	5	Not Graduate	Yes

	loan_amount	loan_term	cibil_score	residential_assets_value
0	29900000	12	778	2400000
1	12200000	8	417	2700000
2	29700000	20	506	7100000
3	30700000	8	467	18200000
4	24200000	20	382	12400000

	commercial_assets_value	luxury_assets_value	bank_asset_value
0	17600000	22700000	NaN
1	2200000	8800000	3300000.0
2	4500000	33300000	12800000.0
3	3300000	23300000	7900000.0
4	8200000	29400000	5000000.0

```
#check for NaN values
```

```
data.isnull().sum()
```

```
Unnamed: 0      0
id              0
dependents      0
education       0
self_employed   0
income_annual   213
loan_amount     0
loan_term       0
cibil_score     0
residential_assets_value  0
commercial_assets_value  0
luxury_assets_value      0
bank_asset_value    341
loan_status       0
dtype: int64
```

```
# Calculate descriptive statistics for numerical columns
```

```
numerical_descriptive_stats = data.describe()
```

```
numerical_descriptive_stats
```

```
      Unnamed: 0      id  dependents  income_annual
loan_amount \
count  4268.000000  4268.000000  4268.000000  4.055000e+03
4.268000e+03
mean    2133.500000  2134.500000    2.499063  5.049420e+06
1.513004e+07
std     1232.209804  1232.209804    1.695954  2.799091e+06
9.041673e+06
min       0.000000    1.000000    0.000000  2.000000e+05
3.000000e+05
25%     1066.750000  1067.750000    1.000000  2.700000e+06
7.700000e+06
50%     2133.500000  2134.500000    3.000000  5.100000e+06
1.450000e+07
75%     3200.250000  3201.250000    4.000000  7.500000e+06
2.150000e+07
max     4267.000000  4268.000000    5.000000  9.900000e+06
3.950000e+07
```

```
      loan_term  cibil_score  residential_assets_value \
count  4268.000000  4268.000000  4.268000e+03
mean    10.900656   599.934396  7.470197e+06
std     5.709840   172.450571  6.502476e+06
min     2.000000   300.000000 -1.000000e+05
25%     6.000000   453.000000  2.200000e+06
50%    10.000000   600.000000  5.600000e+06
75%    16.000000   748.000000  1.130000e+07
```

max	20.000000	900.000000	2.910000e+07
	commercial_assets_value	luxury_assets_value	bank_asset_value
count	4.268000e+03	4.268000e+03	3.927000e+03
mean	4.971556e+06	1.512149e+07	4.958365e+06
std	4.388236e+06	9.099370e+06	3.262255e+06
min	0.000000e+00	3.000000e+05	0.000000e+00
25%	1.300000e+06	7.500000e+06	2.300000e+06
50%	3.700000e+06	1.460000e+07	4.500000e+06
75%	7.600000e+06	2.170000e+07	7.100000e+06
max	1.940000e+07	3.920000e+07	1.470000e+07

Summary of categorical data

```
categorical_descriptive_stats = data.describe(include=['O'])
categorical_descriptive_stats
```

	education	self_employed	loan_status
count	4268	4268	4268
unique	2	2	2
top	Graduate	Yes	Approved
freq	2143	2150	2655

Grouping by education and self_employed to compare loan status counts

```
education_loan_status = data.groupby(['education',
'loan_status']).size().unstack()
self_employed_loan_status = data.groupby(['self_employed',
'loan_status']).size().unstack()
```

education_loan_status, self_employed_loan_status

(loan_status	Approved	Rejected
education		
Graduate	1338	805
Not Graduate	1317	808,
loan_status	Approved	Rejected
self_employed		
No	1317	801
Yes	1338	812)

Dropping unnecessary columns

```
data = data.drop(columns=['Unnamed: 0', 'id'])
```

```

# Handling missing values
imputer = SimpleImputer(strategy='median')
data[['income_annual', 'bank_asset_value']] =
imputer.fit_transform(data[['income_annual', 'bank_asset_value']])

# Encoding categorical variables
data['education'] = data['education'].map({' Graduate': 1, ' Not
Graduate': 0})
data['self_employed'] = data['self_employed'].map({' Yes': 1, ' No':
0})
data['loan_status'] = data['loan_status'].map({' Approved': 1, '
Rejected': 0})

# Filling Negative Values with mean
mean_non_negative = data[data['residential_assets_value'] >= 0]
['residential_assets_value'].mean()
data.loc[data['residential_assets_value'] < 0,
'residential_assets_value'] = mean_non_negative

data.iloc[:, :-1].describe()

```

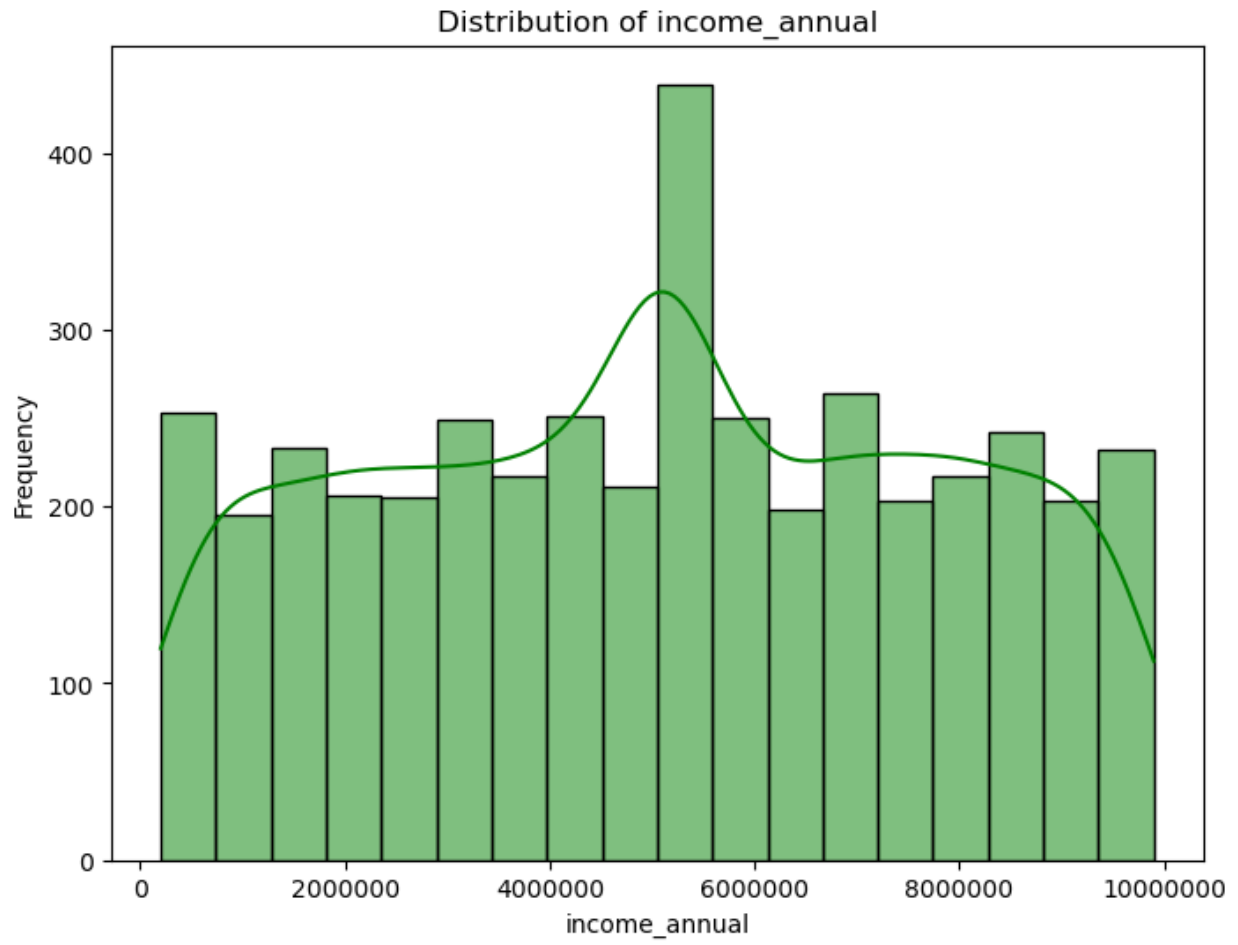
	dependents	education	self_employed	income_annual
loan_amount \				
count	4268.000000	4268.000000	4268.000000	4.268000e+03
mean	2.499063	0.502109	0.503749	5.051945e+06
std	1.695954	0.500054	0.500045	2.728357e+06
min	0.000000	0.000000	0.000000	2.000000e+05
25%	1.000000	0.000000	0.000000	2.800000e+06
50%	3.000000	1.000000	1.000000	5.100000e+06
75%	4.000000	1.000000	1.000000	7.300000e+06
max	5.000000	1.000000	1.000000	9.900000e+06


	loan_term	cibil_score	residential_assets_value \
count	4268.000000	4268.000000	4.268000e+03
mean	10.900656	599.934396	7.520189e+06
std	5.709840	172.450571	6.473304e+06
min	2.000000	300.000000	0.000000e+00
25%	6.000000	453.000000	2.200000e+06
50%	10.000000	600.000000	5.800000e+06
75%	16.000000	748.000000	1.130000e+07
max	20.000000	900.000000	2.910000e+07

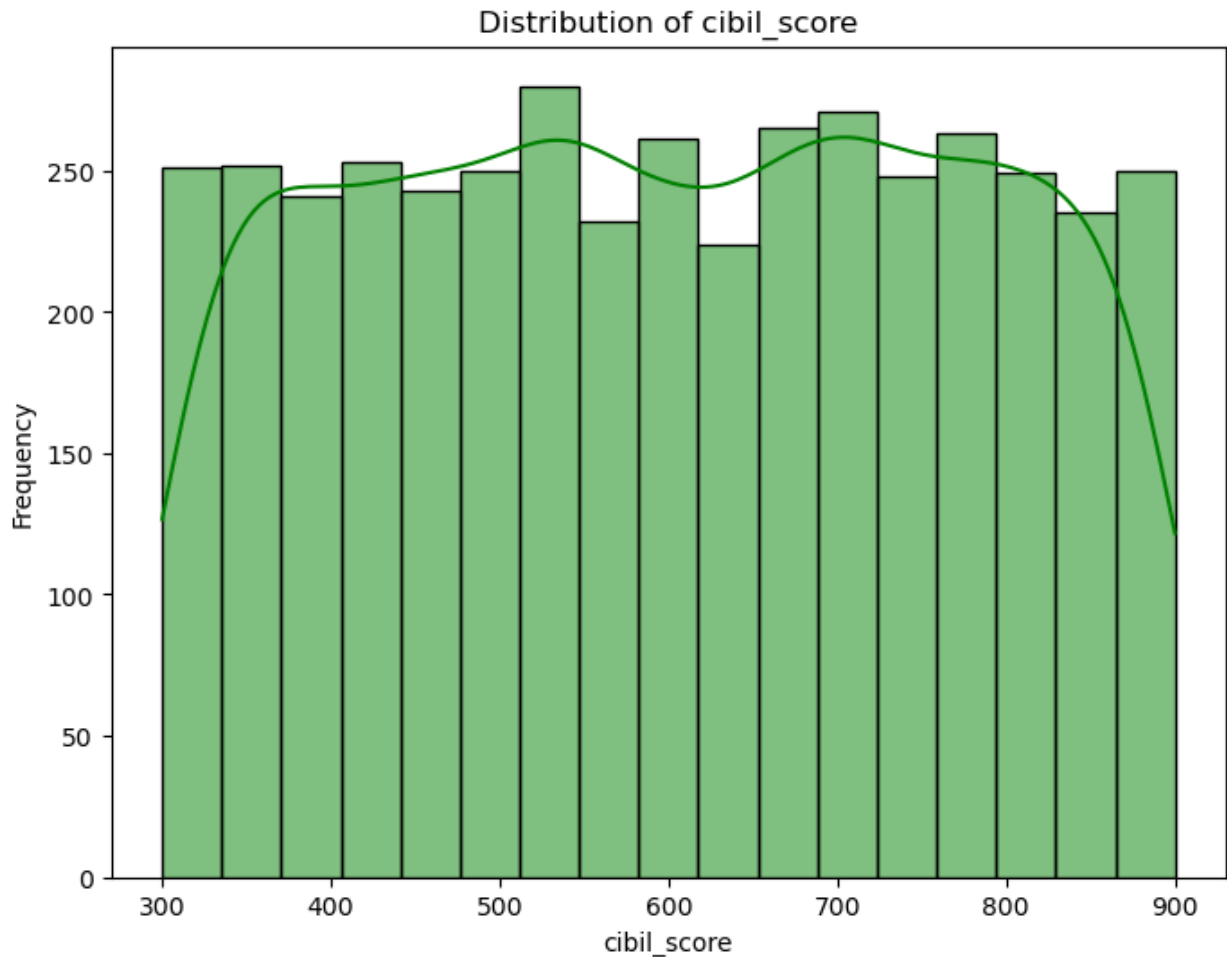
	commercial_assets_value	luxury_assets_value	bank_asset_value
count	4.268000e+03	4.268000e+03	4.268000e+03
mean	4.971556e+06	1.512149e+07	4.921743e+06
std	4.388236e+06	9.099370e+06	3.131656e+06
min	0.000000e+00	3.000000e+05	0.000000e+00
25%	1.300000e+06	7.500000e+06	2.500000e+06
50%	3.700000e+06	1.460000e+07	4.500000e+06
75%	7.600000e+06	2.170000e+07	6.800000e+06
max	1.940000e+07	3.920000e+07	1.470000e+07

```
import seaborn as sns
import matplotlib.pyplot as plt

# Histograms
numerical_vars = ['income_annual', 'loan_amount', 'cibil_score']
for var in numerical_vars:
    plt.figure(figsize=(8, 6))
    ax = sns.histplot(data=data, x=var, kde=True, color='green')
    ax.ticklabel_format(style='plain', axis='x') # Disable scientific
notation
    plt.title(f'Distribution of {var}')
    plt.xlabel(var)
    plt.ylabel('Frequency')
    plt.show()
```




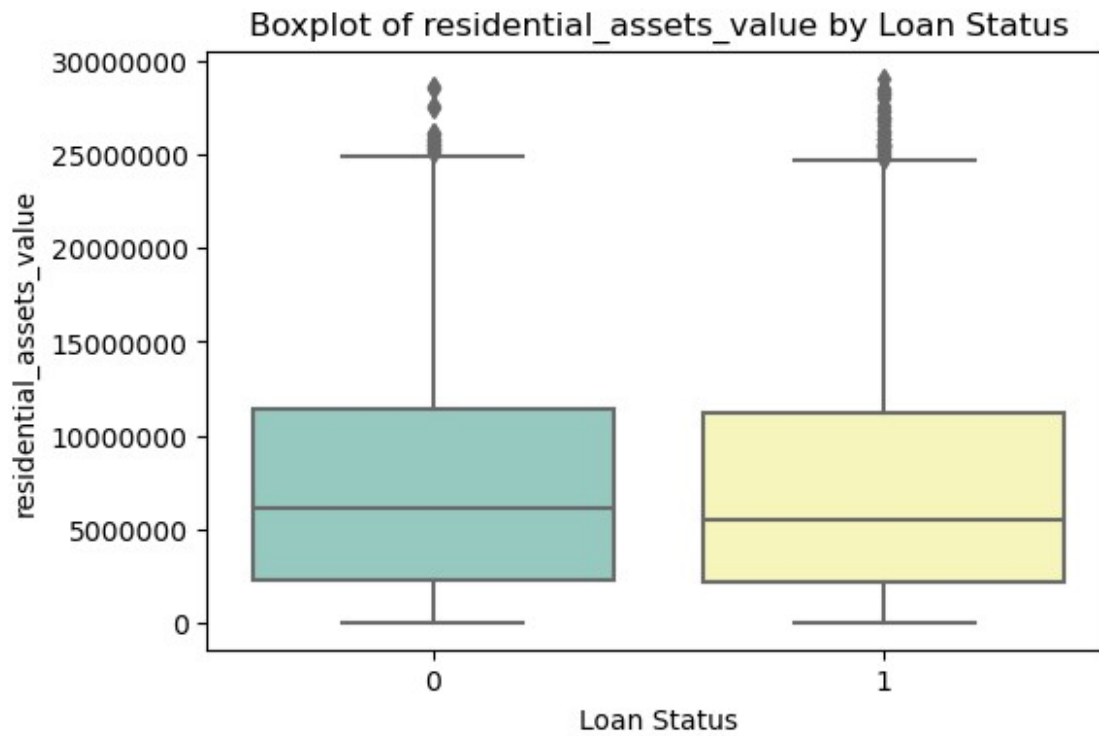






```
import seaborn as sns
import matplotlib.pyplot as plt

# Selecting different numerical variables for box plots
other_numerical_vars = ['residential_assets_value',
                        'commercial_assets_value', 'luxury_assets_value', 'bank_asset_value']

# Boxplots for the selected numerical variables
for var in other_numerical_vars:
    plt.figure(figsize=(6, 4))
    ax = sns.boxplot(x='loan_status', y=var, data=data,
                    palette='Set3')
    ax.ticklabel_format(style='plain', axis='y')
    plt.title(f'Boxplot of {var} by Loan Status')
    plt.xlabel('Loan Status')
    plt.ylabel(var)
    plt.show()
```







```
import seaborn as sns
import matplotlib.pyplot as plt

# Boxplots for all columns
```

```
plt.figure(figsize=(20, 10))
sns.boxplot(data=data, palette='Set2')
plt.title('Boxplot of All Columns')
plt.xticks(rotation=45)
plt.show()
```



```
data.head()
```

	dependents	education	self_employed	income_annual	loan_amount	\
0	2	1	0	9600000.0	29900000	
1	0	0	1	4100000.0	12200000	
2	3	1	0	9100000.0	29700000	
3	3	1	0	8200000.0	30700000	
4	5	0	1	9800000.0	24200000	

	loan_term	cibil_score	residential_assets_value	commercial_assets_value	\
0	12	778	2400000.0	17600000	
1	8	417	2700000.0	2200000	
2	20	506	7100000.0	4500000	
3	8	467	18200000.0	3300000	
4	20	382	12400000.0	8200000	

	luxury_assets_value	bank_asset_value	loan_status
0	22700000	4500000.0	1
1	8800000	3300000.0	0
2	33300000	12800000.0	0
3	23300000	7900000.0	0
4	29400000	5000000.0	0

Calculate the correlation matrix for numerical columns

```
correlation_matrix = data.corr()
```


```
correlation_matrix
```

	dependents	education	self_employed	
income_annual \				
dependents	1.000000	0.002904	0.000557	
0.009312				
education	0.002904	1.000000	-0.022994	
0.013018				
self_employed	0.000557	-0.022994	1.000000	-
0.004497				
income_annual	0.009312	0.013018	-0.004497	
1.000000				
loan_amount	-0.003033	0.010259	0.001831	
0.901647				
loan_term	-0.020145	-0.008381	0.004070	
0.012515				
cibil_score	-0.009991	-0.004659	-0.004856	-
0.019911				
residential_assets_value	0.008234	0.012234	0.007044	
0.617785				
commercial_assets_value	-0.001210	-0.007128	-0.017638	
0.625112				
luxury_assets_value	0.003287	0.011952	0.004950	
0.902909				
bank_asset_value	0.008374	0.008546	-0.001429	
0.796866				
loan_status	-0.017956	0.004737	0.000529	-
0.014084				


	loan_amount	loan_term	cibil_score	\
dependents	-0.003033	-0.020145	-0.009991	
education	0.010259	-0.008381	-0.004659	
self_employed	0.001831	0.004070	-0.004856	
income_annual	0.901647	0.012515	-0.019911	
loan_amount	1.000000	0.008499	-0.017055	
loan_term	0.008499	1.000000	0.007811	
cibil_score	-0.017055	0.007811	1.000000	
residential_assets_value	0.591217	0.007764	-0.018359	
commercial_assets_value	0.602955	-0.005422	-0.003785	
luxury_assets_value	0.860838	0.012581	-0.028656	

bank_asset_value	0.754248	0.012650	-0.015995
loan_status	0.015861	-0.113015	0.770566
residential_assets_value			
commercial_assets_value \			
dependents	0.008234	-	
0.001210			
education	0.012234	-	
0.007128			
self_employed	0.007044	-	
0.017638			
income_annual	0.617785		
0.625112			
loan_amount	0.591217		
0.602955			
loan_term	0.007764	-	
0.005422			
cibil_score	-0.018359	-	
0.003785			
residential_assets_value	1.000000		
0.413127			
commercial_assets_value	0.413127		
1.000000			
luxury_assets_value	0.587511		
0.590825			
bank_asset_value	0.505198		
0.527509			
loan_status	-0.015535		
0.007965			
luxury_assets_value bank_asset_value			
loan_status			
dependents	0.003287	0.008374	-
0.017956			
education	0.011952	0.008546	
0.004737			
self_employed	0.004950	-0.001429	
0.000529			
income_annual	0.902909	0.796866	-
0.014084			
loan_amount	0.860838	0.754248	
0.015861			
loan_term	0.012581	0.012650	-
0.113015			
cibil_score	-0.028656	-0.015995	
0.770566			
residential_assets_value	0.587511	0.505198	-
0.015535			
commercial_assets_value	0.590825	0.527509	

```
# Plotting the heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, fmt=".2f",
            cmap='coolwarm', cbar=True)
plt.title('Heatmap of Correlation Matrix')
plt.show()
```



```
# Creating a violin plot for CIBIL scores by Loan Status
plt.figure(figsize=(10, 6))
sns.violinplot(x='loan_status', y='cibil_score', data=data)
plt.title('Distribution of CIBIL Scores by Loan Status')
plt.xlabel('Loan Status')
plt.ylabel('CIBIL Score')
plt.show()
```



```
# Splitting data into features and target variable
X = data.drop(['loan_status'], axis=1)
Y = data['loan_status']

# Splitting data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, Y,
test_size=0.2, random_state=42)
```

SVM Model

```
from sklearn.preprocessing import StandardScaler

# Create a scaler object
scaler = StandardScaler()

# Fit on training data and transform both training and testing sets
```

```

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

from sklearn.svm import SVC

# Initializing and training the SVM model
svm_model = SVC(kernel='linear', random_state=42)
svm_model.fit(X_train_scaled, y_train)

SVC(kernel='linear', random_state=42)

from sklearn.metrics import accuracy_score, classification_report

# Predict the labels for the test set
y_pred = svm_model.predict(X_test_scaled)

# Calculate the accuracy and classification report
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print("Accuracy:", accuracy)
print("Classification Report:\n", report)

```

Accuracy: 0.9203747072599532

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.89	0.89	303
1	0.94	0.93	0.94	551
accuracy			0.92	854
macro avg	0.91	0.91	0.91	854
weighted avg	0.92	0.92	0.92	854

Hyperparameter Tuning

```

from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC

# Define the model to be tuned
svm_classifier = SVC(random_state=42)

# Define the hyperparameter grid
param_grid = {
    'C': [0.1, 1, 10, 100], # Regularization parameter
    'kernel': ['linear', 'rbf', 'poly'], # Type of SVM kernel to be
used
    'gamma': ['scale', 'auto'], # Kernel coefficient for 'rbf',
'poly'
    'degree': [2, 3, 4] # Degree of the polynomial kernel function

```



```

('poly').
}

# Setup the grid search with cross-validation
svm_grid_search = GridSearchCV(estimator=svm_classifier,
param_grid=param_grid, cv=5,scoring='accuracy', verbose=1)

# Fit the grid search to the data
svm_grid_search.fit(X_train_scaled, y_train)

Fitting 5 folds for each of 72 candidates, totalling 360 fits

GridSearchCV(cv=5, estimator=SVC(random_state=42),
              param_grid={'C': [0.1, 1, 10, 100], 'degree': [2, 3, 4],
                           'gamma': ['scale', 'auto'],
                           'kernel': ['linear', 'rbf', 'poly']},
              scoring='accuracy', verbose=1)

# Print the best parameters and best score
print("Best parameters:", grid_search_svm.best_params_)
print("Best cross-validation score:
{:.2f}".format(svm_grid_search.best_score_))

Best parameters: {'C': 10, 'degree': 2, 'gamma': 'scale', 'kernel':
'rbf'}
Best cross-validation score: 0.94

# Use the best estimator to make predictions on test data
y_pred_svm = svm_grid_search.best_estimator_.predict(X_test_scaled)

# Evaluate the best model on the test set
accuracy_svm = accuracy_score(y_test, y_pred_svm)
print("Test set accuracy: {:.2f}".format(accuracy_svm))

Test set accuracy: 0.94

# Use the best estimator to make predictions on train data
y_pred_svm_train =
grid_search_svm.best_estimator_.predict(X_train_scaled)

# Evaluate the best model on the train set
accuracy_svm_train = accuracy_score(y_train, y_pred_svm_train)
print("Train set accuracy: {:.2f}".format(accuracy_svm_train))

Train set accuracy: 0.98

from sklearn.metrics import precision_score, recall_score, f1_score

# Calculate precision, recall, and F1 score
precision_svm = precision_score(y_test, y_pred_svm)
recall_svm = recall_score(y_test, y_pred_svm)

```

```
f1_svm = f1_score(y_test, y_pred_svm)
```

```
print("Precision:", precision_svm)
```

```
print("Recall:", recall_svm)
```

```
print("F1 Score:", f1_svm)
```

```
Precision: 0.9579524680073126
```

```
Recall: 0.9509981851179673
```

```
F1 Score: 0.9544626593806921
```

Multilayer Perceptron

```
from sklearn.neural_network import MLPClassifier
```

```
from sklearn.metrics import accuracy_score, classification_report
```

```
# Initialize the MLPClassifier
```

```
mlp = MLPClassifier(hidden_layer_sizes=(100, 50), # This is the size of the hidden layer(s)
```

```
                    activation='relu', # Activation function
```

```
                    solver='adam', # The solver for weight
```

```
                    max_iter=300, # Maximum number of
```

```
                    random_state=42)
```

```
# Fit the model to the scaled training data
```

```
mlp.fit(X_train_scaled, y_train)
```

```
MLPClassifier(hidden_layer_sizes=(100, 50), max_iter=300,  
              random_state=42)
```

```
# Make predictions with the model
```

```
y_pred = mlp.predict(X_test_scaled)
```

```
# Calculate the accuracy and print classification report
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
report = classification_report(y_test, y_pred)
```

```
print("Accuracy:", accuracy)
```

```
print("Classification Report:\n", report)
```

```
Accuracy: 0.9601873536299765
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.94	0.95	0.94	303
1	0.97	0.97	0.97	551
accuracy			0.96	854
macro avg	0.96	0.96	0.96	854

weighted avg	0.96	0.96	0.96	854
--------------	------	------	------	-----

Hyperparameter Tuning

```
from sklearn.model_selection import GridSearchCV

# Define the MLP model
mlp = MLPClassifier(max_iter=1000, random_state=42) # Increase
max_iter for convergence

# Define the parameter grid to search
parameter_space = {
    'hidden_layer_sizes': [(50,), (100,), (50, 50), (100, 100)],
    'activation': ['tanh', 'relu'],
    'solver': ['adam'], # 'sgd' removed as it tends to perform worse
    'alpha': [0.0001, 0.05],
    'learning_rate_init': [0.001, 0.01], # Added learning_rate_init
    for learning rate adjustment
}

# Set up GridSearchCV
mlp_grid_search = GridSearchCV(mlp, parameter_space, n_jobs=-1,
cv=3, scoring='accuracy', verbose=2)
```

```
# Fit GridSearchCV
mlp_grid_search.fit(X_train_scaled, y_train)
```

Fitting 3 folds for each of 32 candidates, totalling 96 fits

```
/opt/anaconda3/lib/python3.11/site-packages/sklearn/neural_network/
_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic
Optimizer: Maximum iterations (1000) reached and the optimization
hasn't converged yet.
  warnings.warn(
/opt/anaconda3/lib/python3.11/site-packages/sklearn/neural_network/
_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic
Optimizer: Maximum iterations (1000) reached and the optimization
hasn't converged yet.
  warnings.warn(
/opt/anaconda3/lib/python3.11/site-packages/sklearn/neural_network/
_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic
Optimizer: Maximum iterations (1000) reached and the optimization
hasn't converged yet.
  warnings.warn(
```

```
GridSearchCV(cv=3, estimator=MLPClassifier(max_iter=1000,
random_state=42),
            n_jobs=-1,
            param_grid={'activation': ['tanh', 'relu'],
                        'alpha': [0.0001, 0.05],
```

```

        'hidden_layer_sizes': [(50,), (100,), (50,
50),
                                (100, 100)],
        'learning_rate_init': [0.001, 0.01],
        'solver': ['adam']},
        scoring='accuracy', verbose=2)

# Best parameter set
print('Best parameters found:\n', mlp_grid_search.best_params_)

Best parameters found:
{'activation': 'tanh', 'alpha': 0.05, 'hidden_layer_sizes': (100,
100), 'learning_rate_init': 0.001, 'solver': 'adam'}

means = mlp_grid_search.cv_results_['mean_test_score']
stds = mlp_grid_search.cv_results_['std_test_score']
for mean, std, params in zip(means, stds,
mlp_grid_search.cv_results_['params']):
    print("%0.3f (+/-0.03f) for %r" % (mean, std * 2, params))

0.962 (+/-0.016) for {'activation': 'tanh', 'alpha': 0.0001,
'hidden_layer_sizes': (50,), 'learning_rate_init': 0.001, 'solver':
'adam'}
0.955 (+/-0.010) for {'activation': 'tanh', 'alpha': 0.0001,
'hidden_layer_sizes': (50,), 'learning_rate_init': 0.01, 'solver':
'adam'}
0.958 (+/-0.014) for {'activation': 'tanh', 'alpha': 0.0001,
'hidden_layer_sizes': (100,), 'learning_rate_init': 0.001, 'solver':
'adam'}
0.958 (+/-0.014) for {'activation': 'tanh', 'alpha': 0.0001,
'hidden_layer_sizes': (100,), 'learning_rate_init': 0.01, 'solver':
'adam'}
0.960 (+/-0.015) for {'activation': 'tanh', 'alpha': 0.0001,
'hidden_layer_sizes': (50, 50), 'learning_rate_init': 0.001, 'solver':
'adam'}
0.964 (+/-0.014) for {'activation': 'tanh', 'alpha': 0.0001,
'hidden_layer_sizes': (50, 50), 'learning_rate_init': 0.01, 'solver':
'adam'}
0.965 (+/-0.007) for {'activation': 'tanh', 'alpha': 0.0001,
'hidden_layer_sizes': (100, 100), 'learning_rate_init': 0.001,
'solver': 'adam'}
0.958 (+/-0.033) for {'activation': 'tanh', 'alpha': 0.0001,
'hidden_layer_sizes': (100, 100), 'learning_rate_init': 0.01,
'solver': 'adam'}
0.963 (+/-0.015) for {'activation': 'tanh', 'alpha': 0.05,
'hidden_layer_sizes': (50,), 'learning_rate_init': 0.001, 'solver':
'adam'}
0.962 (+/-0.021) for {'activation': 'tanh', 'alpha': 0.05,
'hidden_layer_sizes': (50,), 'learning_rate_init': 0.01, 'solver':
'adam'}

```

0.964 (+/-0.017) for {'activation': 'tanh', 'alpha': 0.05,
'hidden_layer_sizes': (100,), 'learning_rate_init': 0.001, 'solver':
'adam'}

0.961 (+/-0.014) for {'activation': 'tanh', 'alpha': 0.05,
'hidden_layer_sizes': (100,), 'learning_rate_init': 0.01, 'solver':
'adam'}

0.962 (+/-0.013) for {'activation': 'tanh', 'alpha': 0.05,
'hidden_layer_sizes': (50, 50), 'learning_rate_init': 0.001, 'solver':
'adam'}

0.965 (+/-0.015) for {'activation': 'tanh', 'alpha': 0.05,
'hidden_layer_sizes': (50, 50), 'learning_rate_init': 0.01, 'solver':
'adam'}

0.968 (+/-0.014) for {'activation': 'tanh', 'alpha': 0.05,
'hidden_layer_sizes': (100, 100), 'learning_rate_init': 0.001,
'solver': 'adam'}

0.961 (+/-0.020) for {'activation': 'tanh', 'alpha': 0.05,
'hidden_layer_sizes': (100, 100), 'learning_rate_init': 0.01,
'solver': 'adam'}

0.955 (+/-0.011) for {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (50,), 'learning_rate_init': 0.001, 'solver':
'adam'}

0.951 (+/-0.006) for {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (50,), 'learning_rate_init': 0.01, 'solver':
'adam'}

0.953 (+/-0.021) for {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (100,), 'learning_rate_init': 0.001, 'solver':
'adam'}

0.951 (+/-0.012) for {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (100,), 'learning_rate_init': 0.01, 'solver':
'adam'}

0.953 (+/-0.010) for {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (50, 50), 'learning_rate_init': 0.001, 'solver':
'adam'}

0.959 (+/-0.017) for {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (50, 50), 'learning_rate_init': 0.01, 'solver':
'adam'}

0.953 (+/-0.011) for {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (100, 100), 'learning_rate_init': 0.001,
'solver': 'adam'}

0.951 (+/-0.017) for {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (100, 100), 'learning_rate_init': 0.01,
'solver': 'adam'}

0.960 (+/-0.016) for {'activation': 'relu', 'alpha': 0.05,
'hidden_layer_sizes': (50,), 'learning_rate_init': 0.001, 'solver':
'adam'}

0.958 (+/-0.018) for {'activation': 'relu', 'alpha': 0.05,
'hidden_layer_sizes': (50,), 'learning_rate_init': 0.01, 'solver':
'adam'}

0.955 (+/-0.022) for {'activation': 'relu', 'alpha': 0.05,

```
'hidden_layer_sizes': (100,), 'learning_rate_init': 0.001, 'solver':
'adam'}
0.957 (+/-0.013) for {'activation': 'relu', 'alpha': 0.05,
'hidden_layer_sizes': (100,), 'learning_rate_init': 0.01, 'solver':
'adam'}
0.957 (+/-0.015) for {'activation': 'relu', 'alpha': 0.05,
'hidden_layer_sizes': (50, 50), 'learning_rate_init': 0.001, 'solver':
'adam'}
0.960 (+/-0.010) for {'activation': 'relu', 'alpha': 0.05,
'hidden_layer_sizes': (50, 50), 'learning_rate_init': 0.01, 'solver':
'adam'}
0.959 (+/-0.013) for {'activation': 'relu', 'alpha': 0.05,
'hidden_layer_sizes': (100, 100), 'learning_rate_init': 0.001,
'solver': 'adam'}
0.957 (+/-0.018) for {'activation': 'relu', 'alpha': 0.05,
'hidden_layer_sizes': (100, 100), 'learning_rate_init': 0.01,
'solver': 'adam'}
```

Predict on the test set using the best model

```
y_pred_mlp = mlp_grid_search.best_estimator_.predict(X_test_scaled)
```

```
accuracy_mlp = accuracy_score(y_test, y_pred_mlp)
```

```
accuracy_mlp
```

```
0.9660421545667447
```

Print classification report

```
print(classification_report(y_test, y_pred_mlp))
```

	precision	recall	f1-score	support
0	0.97	0.94	0.95	303
1	0.97	0.98	0.97	551
accuracy			0.97	854
macro avg	0.97	0.96	0.96	854
weighted avg	0.97	0.97	0.97	854

```
from sklearn.metrics import precision_score, recall_score, f1_score
```

Calculate precision, recall, and F1 score

```
precision_mlp = precision_score(y_test, y_pred_mlp)
```

```
recall_mlp = recall_score(y_test, y_pred_mlp)
```

```
f1_mlp = f1_score(y_test, y_pred_mlp)
```

```
print("Precision:", precision_mlp)
```

```
print("Recall:", recall_mlp)
```

```
print("F1 Score:", f1_mlp)
```

Precision: 0.9660714285714286

Recall: 0.9818511796733213

F1 Score: 0.9738973897389739

Predict on the train set using the best model

y_pred_mlp_train =

mlp_grid_search.best_estimator_.predict(X_train_scaled)

accuracy_mlp_train = accuracy_score(y_train, y_pred_mlp_train)

accuracy_mlp_train

0.9824253075571178

from sklearn.metrics import precision_recall_curve

precision_mlp, recall_mlp, _ = precision_recall_curve(y_test,
y_pred_mlp)

precision_svm, recall_svm, _ = precision_recall_curve(y_test,
y_pred_svm)

plt.figure()

plt.plot(recall_mlp, precision_mlp, color='green', lw=2,
label='Precision-Recall curve for MLP')

plt.plot(recall_svm, precision_svm, color='red', lw=2,
label='Precision-Recall curve for SVM')


plt.xlabel('Recall')

plt.ylabel('Precision')

plt.title('Precision-Recall curve')

plt.legend(loc="lower left")

plt.show()




```
from sklearn.metrics import roc_curve, auc

fpr_mlp, tpr_mlp, _ = roc_curve(y_test, y_pred_mlp)
roc_auc_mlp = auc(fpr_mlp, tpr_mlp)

fpr_svm, tpr_svm, _ = roc_curve(y_test, y_pred_svm)
roc_auc_svm = auc(fpr_svm, tpr_svm)

plt.figure()
plt.plot(fpr_mlp, tpr_mlp, color='darkorange', lw=2, label='ROC curve
(area = %0.2f) for MLP' % roc_auc_mlp)
plt.plot(fpr_svm, tpr_svm, color='blue', lw=2, label='ROC curve (area
= %0.2f) for SVM' % roc_auc_svm)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Test Data')
plt.legend(loc="lower right")
plt.show()
```





```
from sklearn.metrics import roc_curve, auc

fpr_mlp, tpr_mlp, _ = roc_curve(y_train, y_pred_mlp_train)
roc_auc_mlp = auc(fpr_mlp, tpr_mlp)

fpr_svm, tpr_svm, _ = roc_curve(y_train, y_pred_svm_train)
roc_auc_svm = auc(fpr_svm, tpr_svm)

plt.figure()
plt.plot(fpr_mlp, tpr_mlp, color='darkorange', lw=2, label='ROC curve
(area = %0.2f) for MLP' % roc_auc_mlp)
plt.plot(fpr_svm, tpr_svm, color='blue', lw=2, label='ROC curve (area
= %0.2f) for SVM' % roc_auc_svm)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Train Data')
plt.legend(loc="lower right")
plt.show()
```

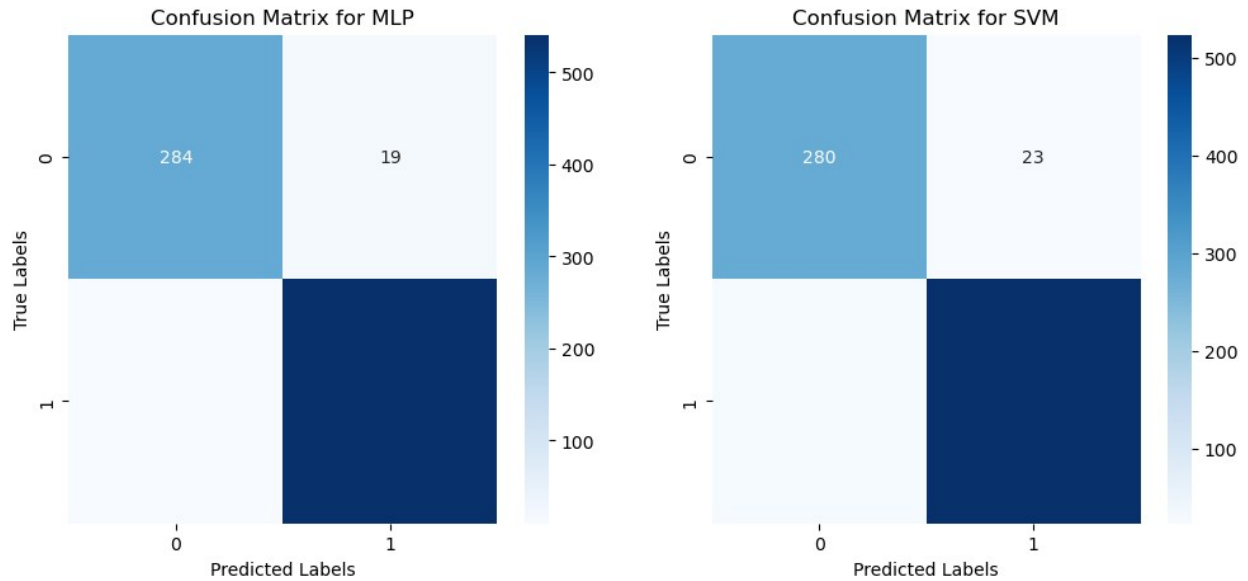


```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

cm_mlp = confusion_matrix(y_test, y_pred_mlp)
cm_svm = confusion_matrix(y_test, y_pred_svm)

fig, ax = plt.subplots(1, 2, figsize=(12, 5))
sns.heatmap(cm_mlp, annot=True, fmt='d', cmap='Blues', ax=ax[0])
ax[0].set_title('Confusion Matrix for MLP')
ax[0].set_xlabel('Predicted Labels')
ax[0].set_ylabel('True Labels')

sns.heatmap(cm_svm, annot=True, fmt='d', cmap='Blues', ax=ax[1])
ax[1].set_title('Confusion Matrix for SVM')
ax[1].set_xlabel('Predicted Labels')
ax[1].set_ylabel('True Labels')
plt.show()
```



```
# We will create a DataFrame with test features and labels
test_data = pd.DataFrame(X_test_scaled)
test_data.reset_index(drop=True, inplace=True)
y_test.reset_index(drop=True, inplace=True)
test_data['True_Label'] = y_test # Add the true labels as a new
column in the DataFrame

# Save the DataFrame to a CSV file
test_data.to_csv('test_data.csv', index=False) # Set index=False not
to save the index

from joblib import dump

# Save the models to a file
dump(mlp_grid_search, 'mlp_grid_search.joblib')
dump(svm_grid_search, 'svm_grid_search.joblib')

['svm_grid_search.joblib']
```